

Can't Tell, or Won't Say?

Separating Evidential and Social Abstention in Multimodal Language Models

Problem: When a multimodal model says “I can’t tell from this image,” the response reads as epistemic humility. But it has two very different causes. *Evidential abstention*: the image genuinely does not support an answer, which is the expected behaviour. *Social abstention*: the image is legible, but the user asserted something confidently and the model hedges to avoid contradicting them. The concerning issue is that abstention is socially rewarded in a way flattery is not; overt agreement invites suspicion, while expressed uncertainty reads as responsible even when it is conflict avoidance. So, sycophancy hiding inside hedging is far harder to detect.




Figure 1: Degradation of one item, normalized to the diagnostic region (the printed code); accuracy anchors come from human calibration (chance = 25%). Colour survives to L3 while the code is lost by L2, so calibration is per category.

Approach: Rather than estimating a model’s uncertainty, we create known levels of uncertainty. We degrade images in calibrated steps and verify with humans where the answer becomes unrecoverable (Fig. 1). Degradation is normalized to the *diagnostic region*, the area determining the answer, and not to raw pixels. We contrast *subtractive* degradation (occlusion; evidence removed, absence visible) with *corruptive* degradation (blur; evidence degraded, nothing announcing absence).

Each stimulus appears under four framings holding content constant and varying only pragmatic force: neutral question, hedged assertion (“*I’m not sure, but is it X?*”), plain assertion, and conviction (“*I’m certain it’s X*”). Assertions are correct on half of trials to separate sycophancy from correctness. Since the image is fixed within a cell, a calibrated model’s abstention should be *invariant* to framing. We quantify departures via **Evidence-Invariance Violation**, $EIV = E_d[\max_a A(d, a) - \min_a A(d, a)]$ for abstention rate A at degradation level d and framing a , and **Pressure-Induced Abstention**, $PIA(d) = A(d, CONV) - A(d, NEUT)$, whose sign separates *dodging* from *capitulating*.

Planned results: Collection is in progress; we state hypotheses in advance. **H1:** abstention varies with framing at matched evidence ($EIV > 0$). **H2:** the effect is largest at *low* degradation—abandoning strong evidence under pressure is a stronger claim than wobbling when evidence is marginal. **H3:** subtractive degradation is better calibrated than corruptive, since absent evidence is legible while degraded evidence is not. A human study ($N \approx 80$) tests **H4:** participants cannot distinguish surface-matched evidential from social abstentions, measured via answer revision.

Context. Abstention benchmarks show models decline far less often than they should, unfixed by scaling [2], but are text-only. Uncertainty-expression studies use scripted hedges in text [3]. Multimodal sycophancy benchmarks [1, 4] are model-side. Closest to us, [1] coarsely lowers image resolution and reports rising flip rates, but without human calibration, an abstention outcome, or graded framing; efforts to separate sycophancy from uncertainty-driven conformity [5] infer model uncertainty rather than control it. We contribute a method for manipulating ground-truth epistemic uncertainty in multimodal settings, a measure separating evidence-driven from pressure-driven abstention, and the first human-subject study of multimodal sycophancy. Prior work asks whether models abstain *enough*; we ask whether their abstention tracks the evidence or the user.

References. [1] Pi et al., On the Sycophancy of Multimodal LLMs, arXiv:2509.16149, 2025. [2] Kirichenko et al., AbstentionBench, arXiv:2506.09038, 2025. [3] Kim et al., “I’m Not Sure, But...”, FAccT 2024, 822–835. [4] Have the VLMs Lost Confidence? (MM-SY), arXiv:2410.11302, ICLR 2025. [5] It’s Not Always Sycophancy (MUSE), arXiv:2605.27288, 2026.